

- `.small-text {font-size: 0.50rem;}`
 - Chapter3: The Relational Database Model
 - Tables and Their Characteristics
 - Relational Database Theory - Functional Dependencies
 - Example 1 of Functional Dependence
 - Example 2 of Functional Dependence
 - Full Functional Dependence
 - Types of Keys
 - Example of Keys
 - Example of a Simple Relational Database
 - Relational Database Keys Comparison
 - Integrity Rules (完整性規則)
 - Integrity Rules - Entity Integrity
 - Integrity Rules - Referential Integrity
 - Illustration of Integrity Rules
 - Relational Database Theory - Relational Algebra
 - Relational Set Operators (SELECT)
 - Relational Set Operators (PROJECT)
 - Relational Set Operators (UNION)
 - Relational Set Operators (INTERSECT)
 - Relational Set Operators (DIFFERENCE)
 - Relational Set Operators (PRODUCT)
 - Relational Set Operators (JOIN)
 - Natural Join
 - Left Outer Join
 - Right Outer Join
 - Relational Set Operators (DIVIDE)
 - Data Dictionary and the System Catalog
 - Relationships within the Relational Database
 - 1:M Relationship
 - 1:1 Relationship
 - M:N Relationship
 - Introduce Composite Entry into M:N Relationship
 - Data Redundancy Revisited
 - Index to Increase Performance
 - Review Questions
-

marp: true theme: default class: invert size: 16:9 paginate: true footer: 國立陽明交通大學
電子與光子學士學位學程 headingDivider: 1 style: | section::after { content: attr(data-
marpit-pagination) '/' attr(data-marpit-pagination-total); }

```
.columns { display: grid; grid-template-columns: repeat(2, minmax(0, 1fr)); gap: 1rem; }  
.columns img { width: 50%; }
```

```
.red-text { color: red; }
```

```
.blue-text { color: LightSkyBlue ;  
}
```

```
.small-text { font-size: 0.50rem; }
```

Chapter3: The Relational Database Model

- Learn about relational database structures like tables, attributes, primary key, foreign key and relationships.
- Learn about how **entity relationship diagrams (ERDs)** can be used to put relational database structures together.

Tables and Their Characteristics

- A table (relation) is a two-dimensional structure composed of rows and columns
- Each row (tuple) represents a single entity occurrence within the entity set
- Each column (attribute) has a distinct column name
- Each intersection of a row and column represents a single data value
- All values in a column must conform to the same data format
- Each column has a specific allowable values known as the attribute domain
- Each table must have primary key (PK) that uniquely identifies each row
- A foreign key (FK) is a primary key of one table that has been placed in another table to illustrate the relationship among tables
- The order of the rows and columns is meaningless

Relational Database Theory - Functional Dependencies

- **Functional dependence** means that the value of one or more attributes determines the value of one or more other attributes
- Definition in relational database theory
 - Given a relation R and attribute sets $X, Y \subseteq R$, X is said to functionally determine Y (written $X \rightarrow Y$) if each X value is associated with precisely one Y value.
 - X is called the **determinant** or **the key**
 - Y is called the **dependent**

Example 1 of Functional Dependence

Student ID	Semester	Lecture	TA
1234	6	Numerical Methods	John
1221	4	Numerical Methods	Smith
1234	6	Visual Computing	Bob
1201	2	Numerical Methods	Peter
1201	2	Physics II	Simon

- StudentID value determinate Semester values. The fact can be expressed by a functional dependency: **StudentID** \rightarrow **Semester**.
- If a row was added (StudentID = 1234 and Semester = 5), then the functional dependency would no longer exist.

Example 2 of Functional Dependence

- $\text{STU_NUM} \rightarrow \text{STU_LNAME}$
- $\text{STU_NUM} \rightarrow (\text{STU_LNAME}, \text{STU_FNAME}, \text{STU_GPA})$
- $(\text{STU_FNAME}, \text{STU_LNAME}, \text{STU_INIT}, \text{STU_PHONE}) \rightarrow (\text{STU_DOB}, \text{STU_HRS}, \text{STU_GPA})$
- $\text{STU_NUM} \rightarrow \text{STU_GPA}$
- $(\text{STU_NUM}, \text{STU_LNAME}) \rightarrow \text{STU_GPA}$


Full Functional Dependence

- Full functional dependence: an attribute is functionally dependent on a composite key but not on any subset of the composite key.
- $\text{STU_NUM} \rightarrow \text{STU_GPA}$ is a full functional dependence
- $(\text{STU_NUM}, \text{STU_LNAME}) \rightarrow \text{STU_GPA}$ is a functional dependence, but NOT a full functional dependence

Types of Keys

- A key consists of attribute(s) that determine other attributes.
 - Determinants in functional dependencies is a key.
 - Eg. invoice number identifies all of the invoice attributes, like invoice date and customer name.
 - Composite key: a key composed of more than one attribute
- Key types
 - Superkey: uniquely identify each entity
 - Candidate key: a minimal Superkey (no extra attributes)
 - PK: choose a candidate key as PK to ensure each row in a table is uniquely identifiable
 - FK: establish relationships among tables

Example of Keys

bg right:50% w:600 keys

- Superkey: STU_NUM , $(\text{STU_NUM}, \text{STU_LNAME})$, $(\text{STU_FNAME}, \text{STU_LNAME}, \text{STU_INIT})$
- Candidate key: STU_NUM , $(\text{STU_FNAME}, \text{STU_LNAME}, \text{STU_INIT})$

- Primary key: STU_NUM
- FK: DEPT_CODE, PROF_NUM

Example of a Simple Relational Database

bg right:60% w:800 relational DB

Relational Database Keys Comparison

Key Type	Definition
Super key	Attribute(s) that uniquely identifies each row
Candidate key	Minimal Superkey without extra attributes
Primary key	Select from candidate keys. Uniquely identify row and cannot be NULL
Foreign key	Is a PK of one table that has been placed in another table to illustrate the relationship among tables

Integrity Rules (完整性規則)

RDBMS rely on integrity rules to ensure data consistency, accuracy, and reliability to prevent errors and enforce business constraints.

- **Entity integrity**
- **Referential Integrity**

Integrity Rules - Entity Integrity

- Rule: Every table must have a PK, and its value cannot be NULL.
- Reason: Ensures that each row in a table is uniquely identifiable, preventing duplicate or missing records.
- Impact w/o it: A db with missing or duplicate keys could lead to data inconsistency.

- Example: When invoice number is a PK, duplicated invoice numbers or empty invoices number is not allowed.

Integrity Rules - Referential Integrity

- Rule: A FK must reference a valid PK in another table.
- Reason: Maintains valid relationships between tables and prevents orphaned records.
- Impact w/o it: If a referenced record is deleted without checking dependencies, it can lead to dangling references.
- Example: A customer might not yet have an assigned sales representative (allow null), but it will be impossible to have an invalid sales representative (must reference).

Illustration of Integrity Rules

 bg right:60% w:800 relational DB

Relational Database Theory - Relational Algebra


- **Relational algebra** defines the theoretical way of manipulating table contents using relational operators.
- Eight main relational operators: SELECT, PROJECT, JOIN, INTERSECT, UNION, DIFFERENCE, PRODUCT, and DIVIDE

Relational Set Operators (SELECT)

SELECT is an operator used to select a subset of rows  bg right:60% w:800 select table

Relational Set Operators (PROJECT)

PROJECT is an operator used to select a subset of columns

 bg right:60% w:800 project table

Relational Set Operators (UNION)

UNION is an operator used to merge two tables into a new table, dropping duplicate rows

bg right:60% w:800 union table

Relational Set Operators (INTERSECT)

INTERSECT is an operator used to yield only the rows that are common to two union-

compatible tables bg right:60% w:800 intersect table


Relational Set Operators (DIFFERENCE)

DIFFERENCE is an operator used to yield all rows from one table that are not found in

another union-compatible table bg right:60% w:800 difference table

Relational Set Operators (PRODUCT)

PRODUCT is an operator used to yield all possible pairs of rows from two tables

bg right:60% w:800 product table

Relational Set Operators (JOIN)

JOIN allows information to be intelligently combined from two or more tables

- **Inner join** – only returns matched records from the tables that are being joined
 - Natural join links tables by selecting only the rows with common values in their common attributes (generally DISCOURAGED in practice)
 - Equijoin – links tables on the basis of an equality condition that compares specified columns of each table
 - Theta join – links tables using an inequality comparison operator
- **Left outer join:** yields all of the rows in the first table, including those that do not have a matching value in the second table
- **Right outer join:** yields all of the rows in the second table, including those that do not have matching values in the first table

Natural Join

PRODUCT -> SELECT -> PROJECT


two tables

select

product


project


Left Outer Join

two tables

left outer join

Right Outer Join

two tables

right outer join

Relational Set Operators (DIVIDE)


The DIVIDE operator is used to answer questions about one set of data being associated with all values of data in another set of data

- Determine which customers (on the left), if any, purchased every product shown in P_CODE table (in the middle).



Data Dictionary and the System Catalog

- **Data dictionary** describes all tables in the DB created by the user and designer
- **System catalog** describes all objects within the database
 - Homonym – same name is used to label different attributes
 - Synonym – different names are used to describe the same attribute.
 - Both homonym and synonym should be avoided whenever possible

bg right:40% w:90% data dictionary

Relationships within the Relational Database

- The one-to-many (1:M) relationship is the norm for relational databases
- In the one-to-one (1:1) relationship, one entity can be related to one and only one other entity and vice versa
- The many-to-many (M:N) relationship can be implemented by creating a new entity in 1:M relationships with the original entities

1:M Relationship

 paint diagram

 paint table

 course diagram

 course table



1:1 Relationship

- 1:1 a professor only chair one department
- 1:M a department employee many professors

 paint diagram

 paint table

M:N Relationship

- A M:N relationship is not supported directly in the relational environment.
- M:N relationship can be implemented by creating a new entity in 1:M relationships with the original entities
- In Fig 3.24, the tables create many data redundancies and relational operation become complex and less efficiency
-  paint diagram
-  paint table

Introduce Composite Entry into M:N Relationship

Table ENROLL is a composite entry (bridge entry, associative entry, link table) to help convert M:N to 1:M

-

Data Redundancy Revisited

- The relational database control of data redundancies through use of foreign keys
- Data redundancy should be controlled except performance and historical data

Index to Increase Performance

- An index is an orderly arrangement to logically access rows in a table
- The index key is the reference point that leads to data location identified by the key
- A table can have many indexes, but each index is associated with only one table
- The index key can have multiple attributes

Review Questions

- What is the integrity rules in RDBMS?
- Describe relational database operators to manipulate relational table contents.
- Describe how to deal with M:N relationship.